

Phobos

A tutorial on the FAT file system

Introduction

This page is intended to provide an introduction to the original File Allocation Table (FAT) file system. This file system was used on all versions of MS-DOS and PC-DOS, and on early versions of Windows; it is still used on floppy disks formatted by Windows and some other systems. Modified versions are also still supported by Windows on hard disks, if required.

The FAT file system is heavily based on the *file map* model in terms of its on-disk layout; that model was around for many years before Microsoft inherited the initial FAT file system from the original writers of DOS (Seattle Computer Products). It is a reasonably simple, reasonably robust file system.

There are three basic variants of the FAT file system, which differ mainly in the construction of the actual file allocation table. Floppy disks and small hard disks usually use the *12-bit* version, which was superseded by the *16-bit* version as hard disks became bigger. This in turn was superseded by the *32-bit* version as disks became bigger still. We shall concentrate on the 16-bit version, since the 12-bit version can be tricky for beginners, and the 32-bit version is more complex than needed for this tutorial.

Overview

Any disk is made up of *surfaces* (one for each head), *tracks* and *sectors*. However, for simplicity, we can consider a disk as a simple storage area made up just of a number of sectors. Further, these sectors are considered to be numbered consecutively, the first being numbered 0, the second numbered 1, etc.; we will not worry about the physical location of any sector on the actual disk. Because we want to emphasise that the location of a sector is irrelevant to the actual disk structure, and because sectors have their own numbers within each track, we shall call these sectors *blocks* from now on; as previously stated, they form a linear, densely numbered list.

All blocks are the same size, 512 bytes, on practically all FAT file systems. However, large disks can have too many blocks for comfort, so blocks are sometimes grouped together in pairs (or fours, or eights, etc...); each such grouping is called an *allocation unit*. The FAT file system actually works in allocation units, not blocks, but for simplicity we shall assume in the description below that each allocation unit contains exactly one block, which means that we can use the terms interchangeably.

A note on numerical values

Hexadecimal numbers are indicated using the convention commonly used in C; that is, a leading 0x. The decimal number 17 would thus be written as 0x11 in hexadecimal notation here.

Values in the FAT file system are either stored in *bytes* (8 bit values, 0-255 unsigned) or

in *words* (pairs of bytes, 16 bit values, 0-65535 unsigned). Note that the first byte of a pair is the least significant byte, and the second byte of a pair is the most significant byte. For example, if the byte at position 3 has a value of 0x15, and the byte at position 4 has a value of 0x74, they together make up a word with value 0x7415 (not 0x1574).

There are occasional 32-bit values (*doublewords*), and these use a similar approach (in this case 4 bytes, with least significant byte stored first).

Lastly, note that individual bits within a byte or word are numbered from the least significant end (right hand end), starting with bit 0.

The disk format

This section describes the *on-disk structure* of a FAT file system; that is, how the various areas of the disk are laid out, and what is stored in them.

Basic layout

All disks using the FAT file system are divided into several areas. The following table summarises the areas in the order that they appear on the disk, starting at block 0:

Area description	Area size
Boot block	1 block
File Allocation Table (may be multiple copies)	Depends on file system size
Disk root directory	Variable (selected when disk is formatted)
File data area	The rest of the disk

The boot block

The boot block occupies just the first block of the disk. It holds a special program (the *bootstrap program*) which is used for loading the operating system into memory. It would thus appear to be fairly irrelevant to this discussion.

However, in the FAT file system it also contains several important data areas which help to describe the rest of the file system. Thus, to understand how a particular disk is laid out, it is necessary first to understand at least part of the contents of the boot block. The relevant areas are shown in the following table, together with their byte offsets from the start of the boot block. We will see, later, which of these are actually important to us.

Offset from start	Length	Description
0x00	3 bytes	Part of the bootstrap program.

0x03	8 bytes	Optional manufacturer description.
0x0b	2 bytes	Number of bytes per block (almost always 512).
0x0d	1 byte	Number of blocks per allocation unit.
0x0e	2 bytes	Number of reserved blocks. This is the number of blocks on the disk that are not actually part of the file system; in most cases this is exactly 1, being the allowance for the boot block.
0x10	1 byte	Number of File Allocation Tables .
0x11	2 bytes	Number of root directory entries (including unused ones).
0x13	2 bytes	Total number of blocks in the entire disk. If the disk size is larger than 65535 blocks (and thus will not fit in these two bytes), this value is set to zero, and the true size is stored at offset 0x20 .
0x15	1 byte	Media Descriptor . This is rarely used, but still exists. .
0x16	2 bytes	The number of blocks occupied by one copy of the File Allocation Table .
0x18	2 bytes	The number of blocks per track. This information is present primarily for the use of the bootstrap program, and need not concern us further here.
0x1a	2 bytes	The number of heads (disk surfaces). This information is present primarily for the use of the bootstrap program, and need not concern us further here.
0x1c	4 bytes	The number of <i>hidden blocks</i> . The use of this is largely historical, and it is nearly always set to 0; thus it can be ignored.
0x20	4 bytes	Total number of blocks in the entire disk (see also offset 0x13).
0x24	2 bytes	Physical drive number. This information is present primarily for the use of the bootstrap program, and need not concern us further here.
0x26	1 byte	Extended Boot Record Signature This information is present primarily for the use of the bootstrap program, and need not concern us further here.
0x27	4 bytes	Volume Serial Number. Unique number used for identification of a particular disk.
0x2b	11 bytes	Volume Label. This is a string of characters for human-readable identification of the disk (padded with spaces if shorter); it is selected when the disk is formatted.

0x36	8 bytes	File system identifier (padded at the end with spaces if shorter).
0x3e	0x1c0 bytes	The remainder of the bootstrap program.
0x1fe	2 bytes	Boot block 'signature' (0x55 followed by 0xaa).

The Media Descriptor

Historically, the size and type of disk were difficult for the operating system to determine by hardware interrogation alone. A 'magic byte' was thus used to classify disks. This are still present, but rarely used, and its contents are known as the Media Descriptor. Generally, for hard disks, this is set to 0xf0.

The File Allocation Table (FAT)

The FAT occupies one or more blocks immediately following the boot block. Commonly, part of its last block will remain unused, since it is unlikely that the required number of entries will exactly fill a complete number of blocks. If there is a second FAT, this immediately follows the first (but starting in a new block). This is repeated for any further FATs.

Note that multiple FATs are used particularly on floppy disks, because of the higher likelihood of errors when reading the disk. If the FAT is unreadable, files cannot be accessed and another copy of the FAT must be used. On hard disks, there is often only one FAT.

In the case of the 16-bit FAT file system, each entry in the FAT is two bytes in length (i.e. 16 bits). The disk data area is divided into *clusters*, which are the same thing as allocation units, but numbered differently (instead of being numbered from the start of the disk, they are numbered from the start of the disk data area). So, the cluster number is the allocation unit number, minus a constant value which is the size of the areas in between the start of the disk and the start of the data area.

Well, almost. The clusters are numbered starting at 2, not 0! So the above calculation has to have 2 added to it to get the cluster number of a given allocation unit...and a cluster number is converted to an allocation unit number by subtracting 2...!

So, how does the FAT work? Simply, there is one entry in the FAT for every cluster (data area block) on the disk. Entry N relates to cluster N. Clusters 0 and 1 don't exist (because of the 'fiddle by 2' above), and those FAT entries are special. The first byte of the first entry is a copy of the [media descriptor](#) byte, and the second byte is set to 0xff. Both bytes in the second entry are set to 0xff.

What does a normal FAT entry for a cluster contain? It contains the *successor cluster number* - that is, the number of the cluster that follows this one in the file to which the current cluster belongs. The last cluster of a file has the value 0xffff in its FAT entry to indicate that there are no more clusters.

The Root Directory

The root directory contains an entry for each file whose name appears at the *root* (the top level) of the file system. Other directories can appear within the root directory; they are called *subdirectories*. The main difference between the two is that space for the root directory is allocated statically, when the disk is formatted; there is thus a finite upper limit on the number of files that can appear in the root directory.

Subdirectories are just files with special data in them, so they can be as large or small as desired.

The format of all directories is the same. Each entry is 32 bytes (0x20) in size, so a single block can contain 16 of them. The following table shows a summary of a single directory entry; note that the offset is merely from the start of that particular entry, not from the start of the block.

Offset	Length	Description
0x00	8 bytes	Filename
0x08	3 bytes	Filename extension
0x0b	1 byte	File attributes
0x0c	10 bytes	Reserved
0x16	2 bytes	Time created or last updated
0x18	2 bytes	Date created or last updated
0x1a	2 bytes	Starting cluster number for file
0x1c	4 bytes	File size in bytes

The Filename

The eight bytes from offset 0x00 to 0x07 represent the filename. The first byte of the filename indicates its status. Usually, it contains a normal filename character (e.g. 'A'), but there are some special values:

0x00

Filename never used.

0xe5

The filename has been used, but the file has been deleted.

0x05

The first character of the filename is actually 0xe5.

0x2e

The entry is for a directory, not a normal file. If the second byte is also 0x2e, the cluster field contains the cluster number of this directory's parent directory. If the parent directory is the root directory (which is statically allocated and doesn't have a cluster number), cluster number 0x0000 is specified here.

Any other character

This is the first character of a real filename.

If a filename is fewer than eight characters in length, it is padded with space

characters.

The Filename Extension

The three bytes from offset 0x08 to 0x0a indicate the filename extension. There are no special characters. Note that the dot used to separate the filename and the filename extension is implied, and is not actually stored anywhere; it is just used when referring to the file. If the filename extension is fewer than three characters in length, it is padded with space characters.

The File Attributes

The single byte at offset 0x0b contains flags that provide information about the file and its permissions, etc. The flags are single bits, and have meanings as follows. Each bit is given as its numerical value, and these are combined to give the actual attribute value:

0x01	Indicates that the file is read only.
0x02	Indicates a hidden file. Such files can be displayed if it is really required.
0x04	Indicates a system file. These are hidden as well.
0x08	Indicates a special entry containing the disk's volume label, instead of describing a file. This kind of entry appears only in the root directory.
0x10	The entry describes a subdirectory.
0x20	This is the archive flag. This can be set and cleared by the programmer or user, but is always set when the file is modified. It is used by backup programs.
0x40	Not used; must be set to 0.
0x80	Not used; must be set to 0.

The File Time

The two bytes at offsets 0x16 and 0x17 are treated as a 16 bit value; remember that the least significant byte is at offset 0x16. They contain the time when the file was created or last updated. The time is mapped in the bits as follows; the first line indicates the byte's offset, the second line indicates (in decimal) individual bit numbers in the 16 bit value, and the third line indicates what is stored in each bit.

```

<----- 0x17 -----> <----- 0x16 ----->
15 14 13 12 11 10 09 08 07 06 05 04 03 02 01 00
h  h  h  h  h  m  m  m  m  m  m  x  x  x  x  x

```

where:

hhhhh

indicates the binary number of hours (0-23)
 mmmmmm
 indicates the binary number of minutes (0-59)
 xxxxx
 indicates the binary number of two-second periods (0-29), representing seconds 0 to 58.

The File Date

The two bytes at offsets 0x18 and 0x19 are treated as a 16 bit value; remember that the least significant byte is at offset 0x18. They contain the date when the file was created or last updated. The date is mapped in the bits as follows; the first line indicates the byte's offset, the second line indicates (in decimal) individual bit numbers in the 16 bit value, and the third line indicates what is stored in each bit.

```

<----- 0x19 -----> <----- 0x18 ----->
15 14 13 12 11 10 09 08 07 06 05 04 03 02 01 00
y  y  y  y  y  y  y  m  m  m  m  d  d  d  d

```

where:

yyyyyyy
 indicates the binary year offset from 1980 (0-119), representing the years 1980 to 2099
 mmmm
 indicates the binary month number (1-12)
 dddd
 indicates the binary day number (1-31)

The Starting Cluster Number

The two bytes at offsets 0x1a and 0x1b are treated as a 16 bit value; remember that the least significant byte is at offset 0x1a. The first cluster for data space on the disk is always numbered as 0x0002. This strange arrangement is because the first two entries in the FAT are reserved for other purposes.

The File Size

The four bytes at offsets 0x1c to 0x1f are treated as a 32 bit value; remember that the least significant byte is at offset 0x1c. They hold the actual file size, in bytes.

Worked examples

The best way to understand how to use the above information is to work through some simple examples.

Interpreting the contents of a block

We assume that there is a tool available to display the contents of a block in both

hexadecimal and as ASCII characters. Most such tools will display unusual ASCII characters (e.g. carriage return) as a dot. For example, here is a display of a typical boot block:

Block 0 (0x0000)																	
	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f	
000	eb	3c	90	49	42	4d	2d	37	2e	30	20	00	02	01	01	00	.<.IBM-7.0
010	01	40	00	a1	13	f8	14	00	0a	00	01	00	00	00	00	00	..@.....
020	00	00	00	00	00	00	29	2a	65	bc	00	43	4f	38	38	33)*e..C0883
030	2d	41	32	20	20	20	46	41	54	31	36	20	20	20	fa	31	-A2 FAT16 .1
040	c0	8e	d0	bc	00	7c	fb	8e	d8	e8	00	00	5e	83	c6	19^...
050	bb	07	00	fc	ac	84	c0	74	06	b4	0e	cd	10	eb	f5	30t.....0
060	e4	cd	16	cd	19	0d	0a	4e	6f	6e	2d	73	79	73	74	65Non-syste
070	6d	20	64	69	73	6b	0d	0a	50	72	65	73	73	20	61	6e	m disk..Press an
080	79	20	6b	65	79	20	74	6f	20	72	65	62	6f	6f	74	0d	y key to reboot.
090	0a	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0a0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0d0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0f0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
100	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
110	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
120	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
130	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
140	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
150	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
160	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
170	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
180	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
190	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1a0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1d0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1f0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	55	aaU.

As an illustration, one field in the boot block has been highlighted in red (the highlight appears twice, once for the hexadecimal representation and once for the ASCII representation). The numbers down the left hand side are the offsets (from the start of the block) of the first byte on that row, and the first row of digits along the top are the offset of each byte within the row. We can thus easily see that the highlighted area starts at offset 0x36.

The area in question is (look back at the boot block layout) the file system type, in this case FAT16. To save us looking up each byte in a table of ASCII characters, we can simply consult the equivalent representation on the right hand side. 0x46 represents F, 0x41 represents A, and so on.

Example 1 - find the root directory

To find the root directory, we need to examine the file system data in the boot block. So, let's look again at the boot block of our example disk:

Block 0 (0x0000)																
	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f
000	eb	3c	90	49	42	4d	2d	37	2e	30	20	00	02	01	01	00
010	01	40	00	a1	13	f8	14	00	0a	00	01	00	00	00	00	00
020	00	00	00	00	00	00	29	2a	65	bc	00	43	4f	38	38	33
030	2d	41	32	20	20	20	46	41	54	31	36	20	20	20	fa	31
040	c0	8e	d0	bc	00	7c	fb	8e	d8	e8	00	00	5e	83	c6	19
050	bb	07	00	fc	ac	84	c0	74	06	b4	0e	cd	10	eb	f5	30
060	e4	cd	16	cd	19	0d	0a	4e	6f	6e	2d	73	79	73	74	65
070	6d	20	64	69	73	6b	0d	0a	50	72	65	73	73	20	61	6e
080	79	20	6b	65	79	20	74	6f	20	72	65	62	6f	6f	74	0d
090	0a	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0a0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0d0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0f0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
100	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
110	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
120	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
130	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
140	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
150	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
160	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
170	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
180	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
190	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1a0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1d0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1f0	00	00	00	00	00	00	00	00	00	00	00	00	00	55	aa	00

We know that the root directory appears immediately after the last copy of the FAT. So what we need to find out is the size of the FAT, and how many copies there are. We also need to know the size of anything else that appears before the FAT(s); there is just the single block of the boot block. So, the number of blocks that appear before the root directory is given by:

$$(\text{size of FAT}) * (\text{number of FATs}) + 1$$

All we need to do, then, is discover these values. First, we know that the number of FATs is stored at offset 0x10 (highlighted in green above); this tells us that there is just one FAT. Next, we need to know the size of a FAT; this is at offsets 0x16 and 0x17,

where we find 0x14 and 0x00 respectively (highlighted in red above). Remember that these two bytes together make up a 16 bit value, with the least significant byte stored first; in other words, the value is 0x0014 (in decimal, 20). So, the total number of blocks that precede the root directory is given by:

$$0x0014 * 1 + 1 \Rightarrow 0x0015 \quad (\text{decimal } 21)$$

We should thus find the root directory in block 0x15, so let's look at it...

Block 21 (0x0015)																
	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f
000	43	4f	38	38	33	2d	41	32	20	20	20	28	00	00	00	00
010	00	00	00	00	00	00	91	9e	65	39	00	00	00	00	00	00
020	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
030	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
040	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
050	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
060	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
070	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
080	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
090	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0a0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0d0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0f0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
100	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
110	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
120	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
130	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
140	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
150	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
160	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
170	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
180	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
190	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1a0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1d0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1f0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00

It seems to have something occupying the first 0x20 bytes, and it's...a directory entry! We won't go into detail here, but detailed examination of those bytes would show that it's the special entry for the disk label. There don't appear to be any more entries in this directory.

Example 2 - find the attributes of a file

In this example, the file FOOBAR.TXT has been created on the same disk, and it appears in the root directory. We wish to find out which attribute flags are set on the file.

First, we need to find the root directory; we have already done this in example 1. Let's take a look at it after FOOBAR.TXT has been created:

Block 21 (0x0015)																	
	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f	
000	43	4f	38	38	33	2d	41	32	20	20	20	28	00	00	00	00	C0883-A2 (....
010	00	00	00	00	00	00	91	9e	65	39	00	00	00	00	00	00e9.....
020	46	4f	4f	42	41	52	20	20	54	58	54	21	00	a3	91	9e	FOOBAR TXT!....
030	65	39	65	39	00	00	91	9e	65	39	c6	10	1a	00	00	00	e9e9....e9.....
040	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
050	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
060	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
070	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
080	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
090	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0a0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0d0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0f0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
100	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
110	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
120	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
130	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
140	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
150	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
160	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
170	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
180	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
190	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1a0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1d0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1f0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00

We can see fairly easily that the second directory entry (the one at offset 0x20) is that for FOOBAR.TXT. Remember that the dot between the filename and the filename extension is not actually stored, but is implied. We see the filename (highlighted in red) and the filename extension (highlighted in blue). We know that the attribute byte appears at offset 0x0b, and it is highlighted in green here.

The value of the attribute byte is 0x21. We can express this in binary as:

0 0 1 0 0 0 0 1

Taking each of the bits separately, and making a hexadecimal number out of them, we get:

```
0 0 1 0   0 0 0 0 => 0x20
0 0 0 0   0 0 0 1 => 0x01
```

Our [table of attribute values](#) shows that 0x20 means that the 'archive flag' is set, and 0x01 indicates that the file is read-only.

Example 3 - find the date of a file

Here, we want the date attached to a particular file (only one date is kept, which is the date of creation or last modification). The file in question is FOOBAR.TXT again.

Let's look once more at the root directory; we have already done this in example 2, and indeed we already know that FOOBAR.TXT has a directory entry at offset 0x20:

Block 21 (0x0015)																	
	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f	
000	43	4f	38	38	33	2d	41	32	20	20	20	28	00	00	00	00	C0883-A2 (....
010	00	00	00	00	00	00	91	9e	65	39	00	00	00	00	00	00e9.....
020	46	4f	4f	42	41	52	20	20	54	58	54	21	00	a3	91	9e	FOOBAR TXT!....
030	65	39	65	39	00	00	91	9e	65	39	c6	10	1a	00	00	00	e9e9.....e9.....
040	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
050	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
060	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
070	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
080	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
090	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0a0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0d0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0f0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
100	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
110	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
120	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
130	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
140	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
150	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
160	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
170	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
180	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
190	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1a0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1d0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1f0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00

This time we are interested in the file date, and we know from our [root directory layout](#) that this is at offset 0x18 within each directory entry. Thus, the date for FOOBAR.TXT is at offset 0x20+0x18, or 0x38 (highlighted in red above). Once again, this is a 16 bit value with the least significant byte stored first. The bytes are 0x65 and 0x39 respectively, so reversing these and putting them together gives a value of 0x3965.

Now all we have to do is analyse the components of this value. An easy way is first to convert it to binary, and this is even easier if we take it one hexadecimal digit at a time:

3	9	6	5
V	V	V	V

0 0 1 1 1 0 0 1 0 1 1 0 0 1 0 1

Let's push all the digits together:

0 0 1 1 1 0 0 1 0 1 1 0 0 1 0 1

Now we can split them again on boundaries corresponding to the individual components of the date, as defined in the [file date format](#). Then we convert each part back to decimal:

0 0 1 1 1 0 0	1 0 1 1	0 0 1 0 1
V	V	V
28	11	5
(year)	(month)	(day)

Remember that the year is based at 1980, so if we add 1980 to 28, we get 2008. The entire date is thus the 5th of November 2008.

Example 4 - find the data blocks for a file

Here, we wish to find out the numbers of the blocks containing data for a particular file which has now been added to the disk. The name of the file is NETWORK.VRS.

Once again, we find the root directory. Here are its latest contents, after NETWORK.VRS has been created:

Block 21 (0x0015)																	
	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f	
000	43	4f	38	38	33	2d	41	32	20	20	20	28	00	00	00	00	C0883-A2 (....
010	00	00	00	00	00	00	91	9e	65	39	00	00	00	00	00	00e9.....
020	46	4f	4f	42	41	52	20	20	54	58	54	21	00	a3	91	9e	FOOBAR TXT!....
030	65	39	65	39	00	00	91	9e	65	39	c6	10	1a	00	00	00	e9e9....e9.....
040	4e	45	54	57	4f	52	4b	20	56	52	53	20	00	b6	91	9e	NETWORK VRS
050	65	39	65	39	00	00	91	9e	65	39	4e 0f	92 06 00 00					e9e9....e9N.....
060	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
070	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
080	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
090	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0a0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0d0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0f0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
100	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
110	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
120	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
130	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
140	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
150	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
160	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
170	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
180	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
190	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1a0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1d0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1f0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00

Note that the third directory entry (starting at offset 0x40) is that for NETWORK.VRS. We know that the starting cluster number for the file data occupies bytes at offsets 0x1a and 0x1b in a particular directory entry; thus the bytes we want are at offsets 0x5a and 0x5b (we just added 0x40, the offset of the start of the entry). These (highlighted in red) contain 0x4e and 0x0f respectively, and, remembering that the first byte is the least significant one, the number we want is 0x0f4e. Incidentally, the next four bytes (highlighted in blue) are the file size, again with the least significant byte first. These are 0x92, 0x06, 0x00, 0x00 respectively, making a value of 0x00000692. This (in decimal) is 1682. So, this file is **1682** bytes long.

Let's review what we know so far...

- The starting cluster of the file is cluster 0x0f4e.
- The root directory starts at block 0x15.
- The first allocation unit starts at the first block after the root directory.

What else do we need to know? We know where the root directory starts, but not where it ends. So we need the size of the root directory, in blocks. Let's look once again at the boot block:

Block 0 (0x0000)																	
	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f	
000	eb	3c	90	49	42	4d	2d	37	2e	30	20	00	02	01	01	00	.<.IBM-7.0
010	01	40	00	a1	13	f8	14	00	0a	00	01	00	00	00	00	00	..@.....
020	00	00	00	00	00	00	29	2a	65	bc	00	43	4f	38	38	33)*e..C0883
030	2d	41	32	20	20	20	46	41	54	31	36	20	20	20	fa	31	-A2 FAT16 .1
040	c0	8e	d0	bc	00	7c	fb	8e	d8	e8	00	00	5e	83	c6	19^...
050	bb	07	00	fc	ac	84	c0	74	06	b4	0e	cd	10	eb	f5	30t.....0
060	e4	cd	16	cd	19	0d	0a	4e	6f	6e	2d	73	79	73	74	65Non-syste
070	6d	20	64	69	73	6b	0d	0a	50	72	65	73	73	20	61	6e	m disk..Press an
080	79	20	6b	65	79	20	74	6f	20	72	65	62	6f	6f	74	0d	y key to reboot.
090	0a	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0a0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0d0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0f0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
100	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
110	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
120	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
130	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
140	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
150	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
160	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
170	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
180	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
190	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1a0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1d0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1f0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	55	aaU.

What we need to find this time is the maximum number of entries in the root directory; this is fixed when the disk is formatted. We know from the [boot block layout](#) that this appears in the two bytes starting at offset 0x11 in the boot block (these are highlighted in red above). These bytes contain 0x40 and 0x00 respectively, so (arranging as usual) this gives us a value of 0x0040 (64 in decimal). So there are 64 root directory entries. We know that one directory entry occupies 32 bytes, so the total space occupied by the root directory is 64*32 bytes, or 2048 bytes. Each block is 512 bytes, so the number of blocks occupied by the root directory is 2048 divided by 512...that is, 4.

So, the root directory starts at block 0x15. Thus the first allocation unit starts at 0x15+4, or 0x19. So, to convert an allocation unit number to a block number, we need to add the constant value 0x19. And to convert a cluster number (which is what

appears in the root directory) to a block number, we need to add 0x17, to allow for that strange offset of 2.

We now know that the first data block of the file is at cluster number 0xf4e (see above). Adding the constant we have discovered, we find that this is block number 0xf4e+0x17, or 0xf65. Let's look at block 0xf65:

Block 3941 (0x0f65)																	
	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f	
000	20	20	20	54	77	61	73	20	74	68	65	20	6e	69	67	68	Twas the nigh
010	74	20	62	65	66	6f	72	65	20	73	74	61	72	74	2d	75	t before start-u
020	70	20	61	6e	64	20	61	6c	6c	20	74	68	72	6f	75	67	p and all throug
030	68	20	74	68	65	20	6e	65	74	2c	0a	20	20	20	20	20	h the net,.
040	6e	6f	74	20	61	20	70	61	63	6b	65	74	20	77	61	73	not a packet was
050	20	6d	6f	76	69	6e	67	3b	20	6e	6f	20	62	69	74	20	moving; no bit
060	6e	6f	72	20	6f	63	74	65	74	2e	0a	20	20	20	54	68	nor octet.. Th
070	65	20	65	6e	67	69	6e	65	65	72	73	20	72	61	74	74	e engineers ratt
080	6c	65	64	20	74	68	65	69	72	20	63	61	72	64	73	20	led their cards
090	69	6e	20	64	65	73	70	61	69	72	2c	0a	20	20	20	20	in despair,.
0a0	20	68	6f	70	69	6e	67	20	61	20	62	61	64	20	63	68	hoping a bad ch
0b0	69	70	20	77	6f	75	6c	64	20	62	6c	6f	77	20	77	69	ip would blow wi
0c0	74	68	20	61	20	66	6c	61	72	65	2e	0a	20	20	20	54	th a flare.. T
0d0	68	65	20	73	61	6c	65	73	6d	65	6e	20	77	65	72	65	he salesmen were
0e0	20	6e	65	73	74	6c	65	64	20	61	6c	6c	20	73	6e	75	nestled all snu
0f0	67	20	69	6e	20	74	68	65	69	72	20	62	65	64	73	2c	g in their beds,
100	0a	20	20	20	20	20	77	68	69	6c	65	20	76	69	73	69	. while visi
110	6f	6e	73	20	6f	66	20	64	61	74	61	20	6e	65	74	73	ons of data nets
120	20	64	61	6e	63	65	64	20	69	6e	20	74	68	65	69	72	danced in their
130	20	68	65	61	64	73	2e	0a	20	20	20	41	6e	64	20	49	heads.. And I
140	20	77	69	74	68	20	6d	79	20	64	61	74	61	73	63	6f	with my datasco
150	70	65	20	74	72	61	63	69	6e	67	73	20	61	6e	64	20	pe tracings and
160	64	75	6d	70	73	0a	20	20	20	20	20	70	72	65	70	61	dumps. prepa
170	72	65	64	20	66	6f	72	20	73	6f	6d	65	20	70	72	65	red for some pre
180	74	74	79	20	62	61	64	20	62	72	75	69	73	65	73	20	tty bad bruises
190	61	6e	64	20	6c	75	6d	70	73	2e	0a	20	20	20	57	68	and lumps.. Wh
1a0	65	6e	20	6f	75	74	20	69	6e	20	74	68	65	20	68	61	en out in the ha
1b0	6c	6c	20	74	68	65	72	65	20	61	72	6f	73	65	20	73	ll there arose s
1c0	75	63	68	20	61	20	63	6c	61	74	74	65	72	2c	0a	20	uch a clatter,.
1d0	20	20	20	20	49	20	73	70	72	61	6e	67	20	66	72	6f	I sprang fro
1e0	6d	20	6d	79	20	64	65	73	6b	20	74	6f	20	73	65	65	m my desk to see
1f0	20	77	68	61	74	20	77	61	73	20	74	68	65	20	6d	61	what was the ma

Well, that certainly looks like the start of a poem! Each line of the text is separated by a special character called *newline*, which has the code 0x0a (decimal 10). The first few of these are highlighted in red.

We have nearly finished. There is obviously more of this file, and for us to find the rest of it, we need to consult the FAT. Recall that the starting *cluster* number of the file (the block we just looked at) is 0xf4e. Each entry in the FAT is two bytes in size, so we'll find the entry for that cluster at offset 0xf4e*2 in the FAT, which is offset 0x1e9c (it's easier to add the value twice than attempt multiplication). We know that one disk block (and

thus one block of the FAT) is 0x200 bytes in size, so we just need to divide 0x1e9c by 0x200. This sounds hard, but it isn't. You can find tools for this, or do it yourself. Let's look at these two numbers in binary:

```
0x0200    => 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0
0x1e9c    => 0 0 0 1 1 1 1 0 1 0 0 1 1 1 0 0
```

The first number is a power of two, so to divide by it we simply shift the second number right - in this case by nine places:

```
0 0 0 0 0 0 0 0 0 0 1 1 1 1    => 0x0f
```

So the entry we want is in block 0x0f of the FAT. The remainder from our division is of course all the bits we lost when we shifted:

```
0 1 0 0 1 1 1 0 0    => 0x9c
```

so this is the byte offset of the entry within the FAT block.

We need to find FAT block 0x0f. We know the FAT starts in block 1 of the disk (see earlier), so block 0x0f of the FAT will be in disk block 0x0f+1, or block 0x10. Let's look at that block:

Block 16 (0x0010)																
	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f
000	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
010	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
020	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
030	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
040	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
050	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
060	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
070	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
080	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
090	00	00	00	00	00	00	00	00	00	00	00	00	4f	0f	50	0f
0a0	51	0f	ff	ff	00	00	00	00	00	00	00	00	00	00	00	00
0b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0d0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0f0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
100	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
110	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
120	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
130	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
140	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
150	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
160	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
170	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
180	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
190	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1a0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1d0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1f0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00

We need to look at the FAT entry (two bytes) at offset 0x9c; this is highlighted in red above, and resolves to the 16 bit value 0x0f4f. This is actually the very next cluster, numerically, from the one we have just looked at (this will not always be the case), so we can apply a bit of common sense and deduce that the second data block of the file appears immediately after the first; thus, the first two blocks are at 0xf65 and 0xf66. Here is block 0xf66:

```

Block 3942 (0x0f66)
  0  1  2  3  4  5  6  7  8  9  a  b  c  d  e  f
000  74 74 65 72 2e 0a 0a 20 20 20 54 68 65 72 65 20 tter...  There
010  73 74 6f 6f 64 20 61 74 20 74 68 65 20 74 68 72 stood at the thr
020  65 73 68 6f 6c 64 20 77 69 74 68 20 50 43 20 69 eshold with PC i
030  6e 20 74 6f 77 2c 0a 20 20 20 20 20 41 6e 20 41 n tow,.  An A
040  52 50 41 4e 45 54 20 68 61 63 6b 65 72 2c 20 61 RPANET hacker, a
050  6c 6c 20 72 65 61 64 79 20 74 6f 20 67 6f 2e 0a ll ready to go..
060  20 20 20 49 20 63 6f 75 6c 64 20 73 65 65 20 66 I could see f
070  72 6f 6d 20 74 68 65 20 63 72 65 61 73 65 73 20 rom the creases
080  74 68 61 74 20 63 6f 76 65 72 65 64 20 68 69 73 that covered his
090  20 62 72 6f 77 2c 0a 20 20 20 20 20 68 65 27 64 brow,.  he'd
0a0  20 63 6f 6e 71 75 65 72 20 74 68 65 20 63 72 69 conquer the cri
0b0  73 69 73 20 63 6f 6e 66 72 6f 6e 74 69 6e 67 20 sis confronting
0c0  68 69 6d 20 6e 6f 77 2e 0a 20 20 20 4d 6f 72 65 him now..  More
0d0  20 72 61 70 69 64 20 74 68 61 6e 20 65 61 67 6c rapid than eagl
0e0  65 73 2c 20 68 65 20 63 68 65 63 6b 65 64 20 65 es, he checked e
0f0  61 63 68 20 61 6c 61 72 6d 0a 20 20 20 20 20 61 ach alarm.  a
100  6e 64 20 73 63 72 75 74 69 6e 69 7a 65 64 20 65 nd scrutinized e
110  61 63 68 20 66 6f 72 20 69 74 73 20 70 6f 74 65 ach for its pote
120  6e 74 69 61 6c 20 68 61 72 6d 2e 0a 0a 20 20 20 ntial harm...
130  4f 6e 20 4c 41 50 42 2c 20 6f 6e 20 4f 53 49 2c On LAPB, on OSI,
140  20 58 2e 32 35 21 0a 20 20 20 20 20 54 43 50 2c X.25!.  TCP,
150  20 53 4e 41 2c 20 56 2e 33 35 21 0a 0a 20 20 20 SNA, V.35!..
160  48 69 73 20 65 79 65 73 20 77 65 72 65 20 61 66 His eyes were af
170  69 72 65 20 77 69 74 68 20 74 68 65 20 73 74 72 ire with the str
180  65 6e 67 74 68 20 6f 66 20 68 69 73 20 67 61 7a ength of his gaz
190  65 3b 0a 20 20 20 20 20 6e 6f 20 62 75 67 20 63 e;.  no bug c
1a0  6f 75 6c 64 20 68 69 64 65 20 6c 6f 6e 67 3b 20 ould hide long;
1b0  6e 6f 74 20 66 6f 72 20 68 6f 75 72 73 20 6f 72 not for hours or
1c0  20 64 61 79 73 2e 0a 20 20 20 41 20 77 69 6e 6b days..  A wink
1d0  20 6f 66 20 68 69 73 20 65 79 65 20 61 6e 64 20 of his eye and
1e0  61 20 74 77 69 74 63 68 20 6f 66 20 68 69 73 20 a twitch of his
1f0  68 65 61 64 2c 0a 20 20 20 20 20 73 6f 6f 6e 20 head,.  soon

```

which certainly looks like the continuation of the poem. If we look at the FAT entry for this new cluster (which, since it's the next block, will also be the next cluster and thus in the next FAT entry), it is highlighted in blue above, and contains the value 0x0f50. This is the very next block and cluster:

```

Block 3943 (0x0f67)
  0  1  2  3  4  5  6  7  8  9  a  b  c  d  e  f
000  67 61 76 65 20 6d 65 20 74 6f 20 6b 6e 6f 77 20 gave me to know
010  49 20 68 61 64 20 6c 69 74 74 6c 65 20 74 6f 20 I had little to
020  64 72 65 61 64 2e 0a 20 20 20 48 65 20 73 70 6f dread.. He spo
030  6b 65 20 6e 6f 74 20 61 20 77 6f 72 64 2c 20 62 ke not a word, b
040  75 74 20 77 65 6e 74 20 73 74 72 61 69 67 68 74 ut went straight
050  20 74 6f 20 68 69 73 20 77 6f 72 6b 2c 0a 20 20 to his work,.
060  20 20 20 66 69 78 69 6e 67 20 61 20 6e 65 74 20 fixing a net
070  74 68 61 74 20 68 61 64 20 67 6f 6e 65 20 70 6c that had gone pl
080  75 6d 62 20 62 65 72 73 65 72 6b 3b 0a 20 20 20 umb berserk;.
090  41 6e 64 20 6c 61 79 69 6e 67 20 61 20 66 69 6e And laying a fin
0a0  67 65 72 20 6f 6e 20 6f 6e 65 20 73 75 73 70 65 ger on one suspe
0b0  63 74 20 6c 69 6e 65 2c 0a 20 20 20 20 20 68 65 ct line,. he
0c0  20 65 6e 74 65 72 65 64 20 61 20 70 61 74 63 68 entered a patch
0d0  20 61 6e 64 20 74 68 65 20 6e 65 74 20 63 61 6d and the net cam
0e0  65 20 75 70 20 66 69 6e 65 21 0a 0a 20 20 20 54 e up fine!.. T
0f0  68 65 20 70 61 63 6b 65 74 73 20 66 6c 6f 77 65 he packets flowe
100  64 20 6e 65 61 74 6c 79 20 61 6e 64 20 70 72 6f d neatly and pro
110  74 6f 63 6f 6c 73 20 6d 61 74 63 68 65 64 3b 0a tocols matched;.
120  20 20 20 20 20 74 68 65 20 68 6f 73 74 73 20 69 the hosts i
130  6e 74 65 72 66 61 63 65 64 20 61 6e 64 20 73 68 nterfaced and sh
140  69 66 74 2d 72 65 67 69 73 74 65 72 73 20 6c 61 ift-registers la
150  74 63 68 65 64 2e 0a 20 20 20 48 65 20 74 65 73 tched.. He tes
160  74 65 64 20 74 68 65 20 73 79 73 74 65 6d 20 66 ted the system f
170  72 6f 6d 20 47 61 74 65 77 61 79 20 74 6f 20 50 rom Gateway to P
180  41 44 3b 0a 20 20 20 20 20 6e 6f 74 20 6f 6e 65 AD;. not one
190  20 62 69 74 20 77 61 73 20 64 72 6f 70 70 65 64 bit was dropped
1a0  3b 20 6e 6f 20 63 68 65 63 6b 73 75 6d 20 77 61 ; no checksum wa
1b0  73 20 62 61 64 2e 0a 20 20 20 41 74 20 6c 61 73 s bad.. At las
1c0  74 20 68 65 20 77 61 73 20 66 69 6e 69 73 68 65 t he was finishe
1d0  64 20 61 6e 64 20 77 65 61 72 69 6c 79 20 73 69 d and wearily si
1e0  67 68 65 64 0a 20 20 20 20 20 61 6e 64 20 74 75 ghed. and tu
1f0  72 6e 65 64 20 74 6f 20 65 78 70 6c 61 69 6e 20 rned to explain

```

We continue this (again, it's the next block and cluster) and we find 0x0f51 as the cluster number (highlighted in green above). Here is that block:

Block 3944 (0x0f68)																	
	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f	
000	77	68	79	20	74	68	65	20	73	79	73	74	65	6d	20	68	why the system h
010	61	64	20	64	69	65	64	2e	0a	20	20	20	49	20	74	77	ad died.. I tw
020	69	73	74	65	64	20	6d	79	20	66	69	6e	67	65	72	73	isted my fingers
030	20	61	6e	64	20	63	6f	75	6e	74	65	64	20	74	6f	20	and counted to
040	74	65	6e	3b	0a	20	20	20	20	20	61	6e	20	6f	66	66	ten;. an off
050	2d	62	79	2d	6f	6e	65	20	69	6e	64	65	78	20	68	61	-by-one index ha
060	64	20	64	6f	6e	65	20	69	74	20	61	67	61	69	6e	2e	d done it again.
070	2e	2e	0a	0a	20	20	20	56	69	6e	74	20	43	65	72	66 Vint Cerf
080	0a	20	20	20	44	65	63	65	6d	62	65	72	20	31	39	38	. December 198
090	35	0a	00	00	00	00	00	00	00	00	00	00	00	00	00	00	5.....
0a0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0d0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0f0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
100	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
110	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
120	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
130	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
140	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
150	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
160	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
170	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
180	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
190	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1a0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1d0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
1f0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00

Lastly, we look at the FAT entry for this block/cluster (highlighted in black). This time the entry is 0xffff, which indicates that there are no more blocks in the file. We have finished!

Conclusion

If you've managed to get this far (and understood it all) you have a good working understanding of the 16-bit FAT file system. You should be able to analyse a disk, and see if it is corrupted. You may even be able to repair it!



This site is copyright © 2012 [Bob Eager](#)

Last updated: 23 May 2012